

In the name of God



University of Tehran

School of Electrical and Computer Engineering

The First Phase of the Final Project

Introduction to Machine Learning

Members:
Arshia Abolghasemi, Mohammad Mahdi Doroudiyan
Kasra Hajiheidari, Shahab Rakhsha

Dey 1403

Contents

Introduction to Voice Authentication 4
Why Voice Authentication is Important 4
Applications of Voice Authentication 5
Open vs. Closed Set Authentication 6
Implementation of Open and Closed Set Authentication 7
Applications of Open and Closed Set Systems 7
Challenges of Voice Authentication and Gender Classification 8
Technical Limitations 8
Security Vulnerabilities 8
Ethical and Privacy Concerns 8
Challenges Specific to Gender Classification 9
Solutions to Challenges in Voice Authentication and Gender Classification 9
Technical Limitations 9
Security Vulnerabilities 10
Ethical and Privacy Concerns 10
Challenges Specifically to Gender Classification 10
Preprocessing of voice data 11
Importance of Preprocessing in Voice Authentication and Gender Classification 11
Why Preprocessing is Important 11
Steps in the Preprocessing Phase 12
1. Noise Reduction 12
2. Normalization 13
3. Windowing 14
Additional Steps in Preprocessing: 15
Feature Extracting Techniques 15
1. Mel Frequency Cepstral Coefficients (MFCC) 15
2. Fast Fourier Transform (FFT) 16
3. Log Mel Spectrogram 17
4. Spectral Centroid 18
5. Chroma Features 19
6. Spectral Contrast 20
7. Zero-Crossing Rate (ZCR) 22
8. Linear Predictive Coding (LPC) 23
9. Perceptual Linear Prediction (PLP) 25
Similarity Learning 27
Similarity Learning Methods 28
Similarity Learning Challenges 29

Different Loss Functions used for Similarity learning 29
Contrastive Loss 29
Triplet Loss 30
Binary Cross-Entropy (BCE) 30
References: 31

Introduction to Voice Authentication

As digital interactions become more common, the need for simple and secure ways to prove who we are has grown. Voice authentication is a modern method that uses the unique features of a person's voice to confirm their identity. Biometrics, which includes traits like fingerprints or facial features, refers to the measurable characteristics that make each person unique. The human voice is a behavioral biometric because it reflects physical and behavioral traits like how our vocal cords, lungs, and speech habits work together to produce sound. Things like pitch, rhythm, and accent make each voice different.

Why Voice Authentication is Important

Voice authentication is a major step forward in making security both easy and reliable. It lets people access secure systems by using their voice instead of needing passwords or physical keys. The process works by recording the user's voice through a microphone, turning it into a digital signal, and comparing it to a stored voice sample to confirm the match. This involves steps like capturing the voice, cleaning up the signal, analyzing features, and deciding if it matches.

Voice authentication is becoming more popular because it's easy to use and fits into everyday technology. Devices like Google Assistant, Siri, and Alexa have made talking to machines a normal part of life. At the same time, many people find it hard to remember all their passwords for the growing number of websites they use. Writing down passwords or using weak ones can lead to security problems. Voice authentication offers a safer and easier option.

From signal acquisition to voice verification, voice signals go through pre-processing, feature extraction, classification, and decision phases, as demonstrated in the figure below:

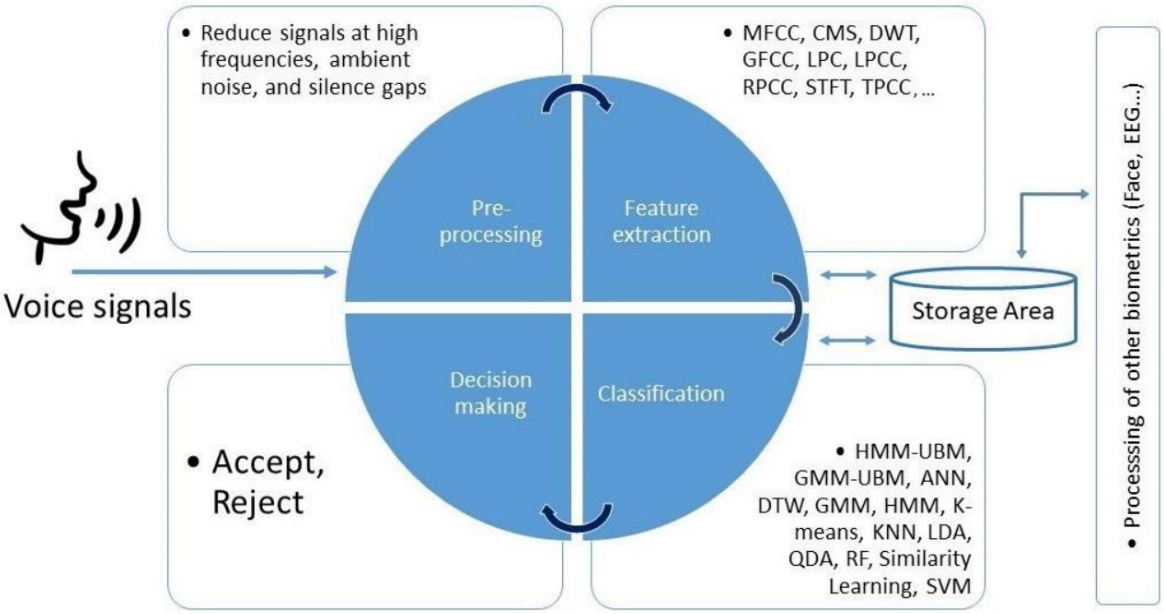


Figure 1: Phases of voice verification

This technology is being used in many areas. In banking, for example, some systems let users perform transactions just by speaking. ATMs are starting to recognize customer voices instead of requiring cards. In healthcare, voice-to-text systems have been used for years to help doctors save time by turning spoken notes into written text. Now, voice authentication also helps patients securely access medical information without needing to visit a hospital, which is especially helpful for older or disabled people.

In conclusion, voice authentication offers a convenient and secure alternative to traditional authentication methods. As technology continues to advance, it is expected that voice authentication will become increasingly prevalent, enhancing both the safety and user-friendliness of various systems.

Applications of Voice Authentication

Voice authentication is used in a wide range of industries and applications:

- **Banking and Financial Services:** Many banks now allow customers to authenticate transactions and access their accounts using voice recognition. ATMs equipped with this technology are becoming more common, offering cardless and secure access.
- **Healthcare:** Voice authentication helps patients access medical records or schedule appointments remotely, especially benefiting elderly or disabled individuals. It also ensures the security of sensitive health information.
- **Customer Service:** Many companies use voice authentication to verify the identity of customers before giving them access to services, reducing wait times and improving security.
- **Smart Home Devices:** Virtual assistants like Siri, Alexa, and Google Assistant use voice recognition to customize user experiences and provide secure access to personal data.
- **Workplace Security:** Businesses use voice authentication for secure access to systems, especially in remote work setups where traditional security measures may be harder to implement.
- **Call center authentication:** Voice biometrics can be used to authenticate customers over the phone, eliminating the need for PIN codes or other forms of identification. This helps to reduce fraud and protect customer data.
- **Government:** Voice biometrics can be used by government agencies for things like border control and passport verification.
- **Law enforcement:** Voice biometrics can be used by law enforcement agencies for things like witness identification and suspect tracking.
- **Education:** Voice biometrics can be used in the education sector for student identification and assessment.
- **Travel:** Voice biometrics can be used in the travel industry for things like airport security and passport control.

- **Telecommunications:** Voice biometrics can be used in the telecommunications industry for customer service and fraud prevention.
- **Gender Classification:** Voice authentication systems can also be used to identify the gender of a speaker based on voice patterns. This application is beneficial in personalization services in various industries.

Open vs. Closed Set Authentication

Voice authentication systems are generally categorized into open-set and closed-set systems, depending on how they handle user data and verification:

- **Closed Set Authentication:** In closed set systems, the system has a predefined group of users, and only these users can be authenticated. For example, a company might set up a system to allow access to employees based on their registered voice samples. These systems are highly secure since they only compare the input voice with a limited number of stored samples.
- **Open Set Authentication:** In open set systems, the system needs to verify whether the input voice belongs to any user, even if the user is not part of a predefined group. This is more challenging because the system must differentiate between enrolled users and unknown voices. For example, a bank's voice authentication system must verify customers from a potentially unlimited user base.

Comparison:

- **Scope of Users:** Closed-set systems work with a fixed group of users, while open-set systems handle a potentially unlimited number of users.
- **Complexity:** Open-set systems are more complex and realistic as they must detect unknown voices and prevent unauthorized access.
- **Use Cases:** Closed set systems are common in controlled environments like workplaces, while open set systems are used in applications like banking and customer service, where new users are frequently enrolled.

Understanding these types helps in designing systems that meet specific security and usability requirements.

Implementation of Open and Closed Set Authentication

Implementing open and closed set authentication systems involves distinct approaches:

- **Closed Set Authentication Implementation:**
  1. **Enrollment:** Collect voice samples from a fixed group of users and store them in the system database.
  2. **Feature Extraction:** Analyze the voice samples to extract unique features such as pitch, tone, and rhythm.
  3. **Matching Algorithm:** Use algorithms to compare input voices with stored samples and determine matches based on similarity thresholds.
  4. **Access Control:** Grant access only if the voice matches one in the predefined database.
- **Open Set Authentication Implementation:**
  1. **Enrollment:** Maintain a database of registered users but also prepare the system to handle unknown inputs.
  2. **Feature Extraction:** Extract unique features from all input voices, both registered and unknown.
  3. **Classification:** Develop a classifier to differentiate between registered and unknown voices. Employ machine learning techniques to enhance detection accuracy.
  4. **Rejection Mechanism:** Include a mechanism to reject voices not matching any registered user with high confidence.

Applications of Open and Closed Set Systems

- **Closed Set Systems:**
  - Used in workplaces or organizations where the user base is fixed, such as employee access to secure areas or systems.
  - Ideal for small-scale applications with a limited number of users.
- **Open Set Systems:**
  - Common in banking and customer service, where the user base is large and dynamic.
  - Suitable for applications like online services and mobile apps where new users are frequently added.
  - Helpful in fraud detection systems, where identifying unauthorized access is critical.

By implementing these methods effectively, voice authentication systems can achieve specific needs across various domains.

Challenges of Voice Authentication and Gender Classification

Voice authentication and gender classification face several challenges that researchers and developers need to address to improve their reliability and security. These challenges can be broadly categorized into technical limitations, security vulnerabilities, and ethical considerations.

Technical Limitations

- **Noisy Environments:** Voice authentication systems often struggle in noisy conditions where background sounds can distort the input voice, reducing accuracy.
- **Variability in Voice:** Factors like aging, illness (e.g., a cold), or stress can change a person's voice, making it difficult for systems to verify users reliably.
- **Computational Resources:** Devices like smart home appliances may lack the processing power and energy required for complex voice authentication algorithms.
- **Standardization Issues:** The absence of universal standards makes it challenging to integrate different voice authentication systems or compare their effectiveness.

Security Vulnerabilities

- **Random Attacks:** Attackers may attempt to access a system without prior knowledge of the target's voice.
- **Mimicry Attacks:** Skilled imitators can replicate a user's voice characteristics, including pitch and tone, to bypass security.
- **Replay Attacks:** Recorded voice samples can be played back to trick authentication systems.
- **Voice Synthesizing Attacks:** Speech synthesis technology can create artificial samples that mimic a user's voice.
- **Hidden Voice Command Attacks:** Inaudible commands embedded in audio signals can manipulate voice-controlled systems without the user's awareness.

To mitigate these risks, techniques such as liveness detection (ensuring the voice comes from a live person), multimodal authentication (combining voice with other biometric or knowledge-based factors), and advanced algorithms are essential.

Ethical and Privacy Concerns

- **Data Privacy:** Collecting and storing voice samples raises concerns about data security and misuse, especially if systems are hacked.
- **User Awareness:** Many users are unaware of the risks associated with sharing biometric data and may inadvertently expose themselves to privacy violations.

Challenges Specific to Gender Classification

- **Cultural and Linguistic Diversity:** Accents and cultural factors can affect how gender is perceived from a voice, making it harder to develop universally accurate systems.
- **Non-Binary Voices:** Gender classification systems often struggle with voices that do not conform to binary categories, limiting inclusivity.
- **Bias in Training Data:** Algorithms trained on biased datasets may reinforce stereotypes or perform poorly with underrepresented groups.

Solutions to Challenges in Voice Authentication and Gender Classification

Addressing the challenges in voice authentication and gender classification involves developing advanced technologies, adopting robust practices, and promoting ethical awareness. Below are solutions categorized by the challenges identified earlier:

Technical Limitations

- **Noisy Environments:**
  - Implement noise-cancellation techniques using advanced signal processing algorithms to isolate the voice signal.
  - Use directional microphones and environmental sound filtering to reduce interference during voice capture.

- **Variability in Voice:**
  - Train systems with diverse datasets that include variations in age, health conditions, and emotional states.
  - Employ adaptive algorithms capable of learning and adjusting to an individual's voice changes over time.
- **Computational Resources:**
  - Optimize algorithms for low-power devices using lightweight models or edge-computing techniques.
  - Use cloud-based processing for devices with limited resources to handle complex computations efficiently.
- **Standardization Issues:**
  - Promote the development of universal protocols and standards for voice authentication to ensure compatibility and interoperability.
  - Collaborate across industries and regulatory bodies to establish benchmarks for performance and security.

Security Vulnerabilities

- **Random and Mimicry Attacks:**
  - Enhance feature extraction techniques to include unique micro-patterns in the voice that are harder to mimic.
  - Combine voice authentication with additional security layers like face recognition or physical tokens for multimodal authentication.
- **Replay and Voice Synthesizing Attacks:**
  - Implement liveness detection methods to ensure the voice input is from a live speaker, not a recording or synthetic generation.
  - Use temporal features like the natural variation in speech patterns to differentiate real-time input from pre-recorded audio.
- **Counterfeit and Hidden Voice Command Attacks:**
  - Develop systems that monitor the context and plausibility of commands before execution.
  - Utilize frequency analysis to detect and block inaudible or hidden commands embedded in malicious audio.

Ethical and Privacy Concerns

- **Data Privacy:**
  - Encrypt all stored and transmitted voice data to protect it from unauthorized access.
  - Allow users control over their data, including options for deletion or restricted use.

- **User Awareness:**
  - Conduct public awareness campaigns to educate users about the risks and best practices for voice authentication.
  - Provide clear terms of service regarding the storage and usage of voice data to ensure informed consent.

Challenges Specifically to Gender Classification

- **Cultural and Linguistic Diversity:**
  - Train models with large and diverse datasets representing different accents, dialects, and speech patterns.
  - Incorporate multilingual capabilities to enhance inclusivity across regions.

- **Bias in Training Data:**
  - Actively collect balanced datasets that include underrepresented groups.
  - Regularly audit algorithms for bias and adjust models to ensure fair performance across demographics.

More Improvements:

- **AI and Deep Learning:**
  Leverage advances in artificial intelligence and deep learning to improve accuracy and adaptability in noisy environments and complex scenarios.

- **Multimodal Systems:**
  Combine voice authentication with other biometrics like facial recognition to address vulnerabilities and enhance security.

- **Continuous Learning:**
  Design systems that update and refine their models based on user feedback and new data, ensuring long-term reliability and effectiveness.

Addressing these challenges with practical and innovative solutions can make voice authentication and gender classification systems more robust, secure, and inclusive.

Preprocessing of voice data

Importance of Preprocessing in Voice Authentication and Gender Classification

Preprocessing is a crucial step in voice authentication and gender classification tasks. It ensures that raw audio data is converted into a clean, standardized, and meaningful format suitable for analysis and machine learning. Without effective preprocessing, the performance and accuracy of models can be significantly compromised.

Why Preprocessing is Important:

- **Noise Reduction:** Real-world audio data often contains background noise. Preprocessing filters out unwanted sounds, improving the quality of the input signal.
- **Feature Extraction:** Raw audio signals are high-dimensional and complex. Preprocessing extracts essential features like Mel-frequency cepstral coefficients (MFCCs) or spectral characteristics, making the data easier to analyze.
- **Standardization:** Audio data may come in different formats, sample rates, or amplitudes. Preprocessing standardizes these aspects to ensure consistency across datasets.
- **Normalization:** Scaling the amplitude of audio signals prevents bias caused by differences in volume or energy levels.
- **Handling Variable Lengths:** Preprocessing segments or pads audio clips to ensure uniform input lengths, which is essential for models requiring fixed-size inputs.

- **Model Efficiency:** Well-preprocessed data can lead to more efficient training and inference processes. It reduces computational load, accelerates convergence during training, and enhances the model's ability to generalize to new, unseen data.
- **Improved Model Performance:** Preprocessing enhances the signal-to-noise ratio which emphasizes relevant features and ensures that the model is provided with high-quality inputs. This leads to improved model performance, accuracy, and robustness.

In voice authentication and gender classification, preprocessing is the foundation that enables advanced algorithms to work efficiently and effectively.

Steps in the Preprocessing Phase

Preprocessing of audio data typically involves multiple steps to prepare the speech signal for analysis. Among these, noise reduction, normalization, and windowing are particularly important.

1. **Noise Reduction**:
    - **Purpose**: Removes background noise, interference, or artifacts that degrade the quality of the audio signal.
    - **Methods**: Three categories of noise removal methods can be distinguished:

      1- Statistical feature-based methods: These are the earliest used methods for speech denoising, Examples of these methods are Wiener filtering, Spectral Subtraction, wavelet transformation denoising, and non-negative matrix factorization.

      2- Denoising Auto-Encoder(DAE)-based methods: DAE was used for the first time in image processing to denoise features for classification. In recent years, the auto-encoder denoising concept has been adopted for extracting clean speech from noisy speech data. DAE is used in various ways for denoising speech, like speaker-aware DAE, time-domain convolutional DAE, and multi-branched encoder DAE.

      3- Deep neural network-based models: These methods utilize Deep Neural Networks (DNNs) to model the non-linear relationship between clean and noisy speech signals. Notably, Google developed a Wavenet-based denoising model that outperforms the traditional Wiener filter. Despite higher computational demands, deep learning methods achieve excellent denoising performance.

    - **Impact**: Enhances the clarity of the speech signal, improving the model's ability to focus on relevant features.

2. **Normalization:**
    - **Purpose**: Balances the amplitude of audio signals to avoid biases caused by varying signal strengths.
    - **Methods**: Scaling audio amplitudes to a fixed range or normalizing based on energy thresholds such as the following functions:

1- Cepstral Mean Normalization (CMN): CMN is the widely used technique for compensating the variability of the speech signal in the cepstral domain. This technique is considered the simplest to implement and the most used normalization technique. The main role of CMN is to remove the convolutional distortions caused by the environmental characteristics of speech (recording device, noises …), at the same time it reduces the effect of speech style. However, the CMN fails to remove convolutional distortions when their time constant approaches the length of the analysis window.

2- Cepstral Variance Normalization (CVN): CVN is a similar technique to CMN, which consists of estimating the variance of each cepstral dimension and then normalizing it into a unity. CVN is not associated with any particular type of distortion. However, it does offer some robustness against speaker variability, additive noise, and acoustic channels.

3- Cepstral Mean and Variance Normalization (CMVN): CVN and CMN are often paired together as the cepstral mean and variance normalization. Therefore, CMVN benefits from both the sample mean and standard deviation, to normalize the cepstral sequence. The use of CMVN positively affects the accuracy and reduces the error rate when using a clean or multi-style acoustic model. Unlike CMN, CMVN shifts and scales the energy term.

4- Histogram Equalization (HEQ): HEQ transforms the cepstral coefficients by normalizing the probability density function and making it equal to the reference probability density function. HEQ is frequently used in Digital Image Processing but has become an important normalization technique for robust speech processing. Compared to CMNV, HEQ doesn't only eliminate the linear effects of noises but also eliminates the non-linear distortions caused by noise.

5- Cepstral Gain Normalization (CGN): CGN is similar to CVN and MVN. However, CVN and MVN are not sufficient to identify the effects of additive noise. The philosophy of CGN is based on the approximate model, which makes this normalization method robust to noise. The application of CGN starts by Subtracting the average of cepstral coefficients which is known as CMN. Then Normalize gain to unity by calculating the maximum

and the minimum sample values of cepstral dimension coefficients. The main benefit of using CGN is the gain of noise-robust performance whatever the environment's noise and without distorting the original speech signal.

6- Quantile-based Cepstral dynamics Normalization (QCN): QCN is a cepstral dynamics normalization technique inspired by CVN, CMN, and CGN techniques, which is mainly used to reduce the sensitivity of the normalization to outliers. QCN uses the quantile estimates for each cepstral dynamic dimension. Compared to the other normalization techniques QCN gives an accurate alignment of the cepstral distribution of samples and proves robustness to distributions shapes changes and reverberations.

7- Wiener-filtering (WF): WF is a normalization technique that is used to remove the additive noise on the basics of wiener gains. WF considers that the speech and noises are not correlated and minimizes the mean square error. Also, there are some weaknesses to be highlighted, WF estimates the power of clean signal and noise before filtering ignoring the non-stationary nature of speech, also it considers gain function as fixed at every frequency.

- **Impact**: Ensures consistent signal representation across the dataset, improving training stability and accuracy.

3. **Windowing:**
   - **Purpose**: Smooth the transition between frames of the speech signal to reduce spectral distortion.
   - **Methods**: There exist several window functions, for spectral performance analysis and digital finite impulse response filter, which can be divided into two groups:

     1- Fixed window functions: rectangular window, Hann window, Hamming window, Blackman window...

     2- Adjustable window functions: Kaiser window, Saramäki window, ultraspherical window, etc.

   - **Impact**: Helps maintain continuity in spectral analysis, improving feature extraction and overall performance.

Additional Steps in Preprocessing:

- **Voice Activity Detection (VAD)**: Identifies and keeps only the voiced parts of speech, eliminating silence and non-speech regions. Allows to meet the limited hardware capabilities and to optimize the CPU consumption.
- **Pre-emphasis**: The frequency components of a speech signal fall in the high frequencies, the pre-emphasis is applied to flatten the magnitude spectrum and balance the high and low-frequency components of the speech signal.
- **Framing**: Since the speech signal has a non-stationary nature, the properties of a signal are not static over time, However, in short time intervals, the speech signal is considered as a stationary signal and easy to process. The Framing step fills this gap by splitting the continuous speech signal into a series of blocks called frames of 20-40ms length to allow block processing of the speech signal.

Feature Extracting Techniques

Feature extraction is a vital step in audio data analysis, as it transforms raw audio signals into a set of representative attributes that models can analyze. These features capture essential information about the audio signal while reducing its dimensionality, improving the efficiency and accuracy of downstream tasks like voice authentication and gender classification. Below are some commonly used techniques:

1. Mel Frequency Cepstral Coefficients (MFCC)

Mel Frequency Cepstral Coefficients (MFCCs) are one of the most widely used methods for feature extraction in audio and speech processing. They represent a compact and robust representation of the spectral properties of sound, which is especially useful for tasks like speech recognition, speaker identification, and music classification.

MFCCs aim to mimic the human auditory system by representing sound on the Mel scale, a perceptual scale of pitches as perceived by human hearing. The process involves transforming an audio signal into a set of features that reflect the relevant properties of the sound while discarding unnecessary information (such as noise or redundant data).

- **Purpose:** Captures the timbral characteristics of speech, mimicking how the human ear perceives sound.
- **Process:**

    1. **Pre-emphasis:** involves emphasizing higher frequencies in the audio signal. This is done using a high-pass filter.
    2. **Framing:** The continuous audio signal is divided into short overlapping frames (typically 20–40 milliseconds long).
    3. **Windowing:** Each frame is multiplied by a window function (such as a Hamming window) to reduce spectral leakage.

4. **Fourier Transform (FFT):** The Fast Fourier Transform (FFT) is applied to convert each frame from the time domain to the frequency domain.
5. **Mel Filter Bank:** The spectrum is passed through a series of Mel filter banks. These filters are triangular and spaced according to the Mel scale.

$$Mel(f) = 2595 \times \log_{10}(1 + \frac{f}{700})$$

6. **Logarithm (Log Energy):** The output of each Mel filter is summed, and the logarithm is applied.
7. **Discrete Cosine Transform (DCT):** The Discrete Cosine Transform (DCT) is applied to log energies to decorrelate features and compress the data into a smaller set of coefficients.
8. **Pass the spectrum through a Mel-scale filter bank to emphasize frequencies critical for human speech perception.**

- **Applications:**
  - **Speech Recognition:** ASR (Automatic Speech Recognition) systems use MFCCs to extract phoneme-level features.
  - **Speaker Identification:** MFCCs capture voice-specific characteristics.
  - **Music Analysis:** Used in music genre classification and melody recognition.
  - **Environmental Sound Classification:** Helps in recognizing specific audio environments (e.g., urban noise, wildlife sounds).

By extracting MFCCs, an audio signal can be represented in a concise and information-rich way, facilitating better performance in machine learning models for audio-related tasks.

2. Fast Fourier Transform (FFT)

The Fast Fourier Transform (FFT) is an efficient algorithm to compute the Discrete Fourier Transform (DFT) and its inverse. It converts a signal from the time domain into the frequency domain, allowing us to analyze the frequency components of the signal. The DFT is computationally expensive, FFT optimizes this process, reducing the complexity and making it much faster for practical applications.

- **Purpose:** Analyze the frequency content of a signal by transforming it from the time domain to the frequency domain. It breaks a complex waveform (signal) into its constituent frequencies, represented as sine and cosine waves.
- **Process:**
  1. Divide the signal into frames.
  2. Apply FFT to each frame to compute the frequency spectrum.
- **Mathematics:** DFT is defined as:

$$X[k] = \sum_{n=0}^{N-1} x[n] \cdot e^{-j\frac{2\pi}{N}kn}$$

- **Applications:**
    - Speech Recognition: Extracting spectral features.
    - Image Processing: Frequency-based filtering.
    - Audio Processing: Analyzing sound frequencies.
    - Signal Analysis: Identifying dominant frequencies in signals.

In essence, FFT is a cornerstone in signal processing for converting signals to their frequency-domain representations efficiently.

3. Log Mel Spectrogram

It is a widely used feature extraction method in audio and speech processing. It combines the principles of a **spectrogram** and the **Mel scale**, followed by applying a logarithmic transformation to better represent how humans perceive sound. The Log Mel Spectrogram is compact and easier to work with compared to raw spectrograms.

- **Key Concepts:**
    1. **Spectrogram**:
       A spectrogram is a visual representation of the spectrum of frequencies in a signal over time. It is obtained by applying a **Short-Time Fourier Transform (STFT)** to the audio signal to convert it into the frequency domain and the result is a matrix.
    2. **Mel Scale**:
       The Mel scale is a perceptual scale that reflects how humans hear sound frequencies. It emphasizes lower frequencies (where humans are more sensitive) and de-emphasizes higher frequencies. Frequencies from the spectrogram are mapped to the Mel scale using a series of **Mel filter banks**.
    3. **Logarithmic Transformation**:
       The Mel spectrogram values are passed through a logarithmic function to compress the dynamic range. This step mimics the human perception of loudness, which is logarithmic.
- **Purpose:** Visualizes the frequency content of a signal over time on a logarithmic scale.
- **Process:**
    1. **Pre-processing**: Pre-emphasize the audio signal to amplify high frequencies. Divide the signal into short overlapping frames (e.g., 20–40 ms).

2. **Fourier Transform**: Apply **STFT** to each frame to get the spectrogram (frequency vs. time).
3. **Mel Filter Bank**: Pass the spectrogram through a set of **Mel filter banks**. Each filter is triangular and focuses on a specific frequency range on the Mel scale.
4. **Sum and Convert**: Sum the energy under each Mel filter to get a Mel spectrogram.
5. **Logarithmic Transformation**: Apply the logarithmic function to the Mel spectrogram values to get the **Log Mel Spectrogram**.

- **Applications:**

1. **Speaker Identification**: Helps in identifying unique voice characteristics.
2. **Music Analysis**: Used in music genre classification and melody recognition.
3. **Speech Recognition:** Commonly used as an input feature in systems like Automatic Speech Recognition (ASR).
4. **Sound Classification**: Useful in tasks like environmental sound recognition.

4. Spectral Centroid

The Spectral Centroid is a key feature in audio signal analysis that indicates the "center of mass" of the spectrum of a sound. It measures where the majority of the energy in a frequency spectrum is concentrated, providing a perceptual indicator of the sound's brightness.

It is mathematically defined as the weighted mean of the frequencies present in the spectrum, with their magnitudes as the weights.

$$Spectral\ Centroid = \frac{\sum_k f_k \cdot |X(f_k)|}{\sum_k |X(f_k)|}$$

f_k: Frequency bin k (in Hz).

$|X(f_k)|$: Magnitude of the spectrum at frequency f_k.

$\sum_k |X(f_k)|$: Total energy of the spectrum.

- **Key Intuition**

A low Spectral Centroid means the energy is concentrated in the lower frequencies, resulting in a "darker" or "bass-heavy" sound. A high Spectral Centroid means the energy is concentrated in the higher frequencies, resulting in

a "brighter" or "sharp" sound. For example, a violin playing a high note has a high spectral centroid and a bass guitar has a low spectral centroid.

- **Purpose:** Represents the "center of mass" of the spectrum, indicating where the majority of the signal's energy is concentrated.
- **Process:**
  - **Pre-processing:** Pre-emphasize and frame the signal. Apply a window function (e.g., Hamming window) to reduce spectral leakage.
  - **Fourier Transform:** Compute the magnitude spectrum using the FFT.
  - Calculate the Centroid: Use the formula above to find the weighted average of the frequencies using their magnitudes as weights.
  - **Output over Time:** For time-varying signals, compute the Spectral Centroid for each frame to observe how the brightness evolves.
- **Applications:**
  - **Audio Analysis:** Measure the perceptual "brightness" of audio.
  - **Music Genre Classification:** Genres with brighter sounds (e.g., electronic or pop) tend to have higher Spectral Centroids compared to bass-heavy genres like jazz or classical.
  - **Speech Processing:** Helps in identifying different voice qualities, such as distinguishing between clear and muffled speech.
  - **Sound Effects:** Used in audio editing and sound design to analyze or manipulate timbre.

By analyzing the Spectral Centroid, we gain insight into the tonal and perceptual characteristics of an audio signal, making it an important feature in speech and audio analysis.

5. Chroma Features

Chroma features (or chroma vectors) are powerful audio features used in music analysis, especially for capturing harmonic and tonal content. They represent the distribution of energy across the 12 distinct pitch classes (or chroma bins) of the musical octave, regardless of the octave in which the pitches occur.

The term chroma refers to the 12 pitch classes in Western music (C, C#, D, D#, E, F, F#, G, G#, A, A#, B). Chroma features map all pitches (frequencies) in an audio signal to one of these 12 chroma bins.

- **Purpose:** Captures pitch-related features by grouping frequencies into 12 pitch classes corresponding to the musical notes.
- **Process:**
  - **Pre-processing:** Apply a Short-Time Fourier Transform (STFT) to the audio signal to obtain the magnitude spectrum (frequency vs. time).

- **Frequency Mapping to Chroma Bins:** Map each frequency component to one of the 12 chroma bins based on its pitch class. The mapping is achieved by considering both the frequency and its corresponding pitch class on the chromatic scale.
- **Aggregation Across Octaves:** Combine the energies of all frequencies that belong to the same pitch class, regardless of the octave.
- **Normalization:** Normalize the resulting chroma vector to ensure that the feature values are comparable across frames.

- **Applications:**
  - **Music Information Retrieval:** Identify the key of a song. Detect chord progressions or harmonies. Perform melody extraction.
  - **Genre Classification:** Different genres exhibit distinct harmonic patterns that can be captured using chroma features.
  - **Cover Song Identification:** Compare harmonic structures of songs to identify covers, since the harmonic content is preserved even if the instrumentation differs.
  - **Audio Fingerprinting:** Match songs or melodies by their harmonic content.
- **Comparison to Other Features**

  Unlike the Mel spectrogram, which focuses on frequency and amplitude, chroma features focus purely on the harmonic content (pitch classes) and ignore octave information.

  In contrast to MFCCs, which are useful for speech processing, chroma features are more suited for music and tonal analysis.

Chroma features are a compact and musically meaningful representation of audio signals, particularly useful for harmonic and tonal analysis. They are widely applied in music information retrieval, genre classification, and audio analysis tasks.

6. Spectral Contrast

Spectral Contrast is a feature used in audio and speech analysis that measures the difference (or contrast) between the peaks and valleys in the spectrum of an audio signal. It captures the dynamic range of spectral energy across different frequency bands, making it particularly useful for distinguishing different audio textures, timbres, and sound types.

- Peaks: High-energy regions in the spectrum (e.g., harmonics, formants).
- Valleys: Low-energy regions in the spectrum (e.g., noise or silence between harmonics).

Spectral contrast measures the difference in energy between these peaks and valleys across multiple frequency bands.

- **Key Intuition:** Sounds with high spectral contrast tend to have sharp differences between peaks and valleys, like in percussive sounds or music with pronounced harmonics. Sounds with low spectral contrast have smoother spectra, like in noise or heavily reverberated signals. For example, a flute has low spectral contrast because its harmonics are smooth and close in energy but a drum hit has high spectral contrast because of sharp transients and pronounced frequency differences.

- **Purpose:** Measures the difference in amplitude between peaks and valleys in the spectrum, providing information about the texture and timbre of a sound.
- **Process:**
    - Pre-processing: Apply an STFT to the signal to obtain its magnitude spectrum over time.
    - Divide the Spectrum into Frequency Bands: The spectrum is divided into several logarithmically spaced frequency bands, typically mimicking the human auditory system.
    - Identify Peaks and Valleys: For each frequency band, identify the local maxima (peaks) and local minima (valleys) in the magnitude spectrum.
    - Compute the Contrast: Compute the difference between the average energy of the peaks and the average energy of the valleys within each band.
    - Output: The result is a spectral contrast vector, where each element corresponds to the contrast value of a specific frequency band.
- **Applications:**
    - Music Genre Classification: Spectral contrast is effective at distinguishing genres with different timbral qualities.
    - Speech Processing: Used for speaker identification or emotion recognition, as different voices and emotions produce varying spectral contrasts.
    - Environmental Sound Classification: Helps distinguish between sharp, transient sounds (e.g., sirens) and smooth, continuous sounds (e.g., wind noise).
    - Music Analysis: Useful for analyzing the timbre or tonal characteristics of musical instruments.
- **Key Differences from Other Features**

- Spectral Centroid: Focuses on the "center of mass" of the spectrum, while spectral contrast measures the range (difference) of energy within bands.
- MFCCs: Focus on the overall spectral shape, while spectral contrast captures fine-grained variations (peaks and valleys).

Spectral contrast is a feature that captures the dynamic range of energy in different frequency bands, providing insights into the timbral and textural properties of sound. Its ability to capture subtle spectral variations makes it an important complement to other spectral features.

7. Zero-Crossing Rate (ZCR)

ZCR is a basic and widely used feature in audio signal processing. It measures how frequently the signal waveform crosses the horizontal axis (i.e., changes its sign from positive to negative or vice versa) within a given time frame. This feature is particularly useful for distinguishing between voiced and unvoiced speech, as well as for music and audio classification. It is computationally simple and also can be used in real-time applications.

It is mathematically defined as:

$$ZCR = \frac{1}{N-1}\sum_{n=1}^{N-1} 1[x(n) \cdot x(n-1) < 0]$$

Where:

- N: Number of samples in the frame.
- $x(n)$: The audio signal at the n-th sample.
- $1[\cdot]$: Indicator function that outputs 1 if the condition is true, otherwise 0.
- $x(n) \cdot x(n-1) < 0$: Checks if the signal changes sign between two consecutive samples.

- High ZCR: Indicates rapid oscillations in the waveform (e.g., noise or unvoiced speech like "s", "f", or percussive sounds in music).
- Low ZCR: Indicates a smoother waveform with fewer sign changes (e.g., voiced speech like vowels or sustained musical notes).

- **Purpose:** Indicates the rate at which the signal changes sign (crosses the zero axis).
- **Process:**
  - Divide the Signal into Frames: Break the audio signal into short frames (e.g., 20–40 ms) for short-term analysis.

- Count Sign Changes: For each frame, count the number of times the signal crosses zero (i.e., changes from positive to negative or vice versa).
- Normalize: Normalize by the total number of samples in the frame to compute the ZCR.

- **Applications:**
  - Speech Processing: Can be used to detect voiced or unvoiced sounds. Voiced sounds (e.g., vowels) have low ZCR, and unvoiced sounds (e.g., "s", "f") have high ZCR. It is also useful in speech segmentation and recognition systems.
  - Music Analysis: Helps differentiate between smooth, harmonic sounds (e.g., flute) and noisy, percussive sounds (e.g., drums or cymbals). Often used in beat detection.
  - Genre Classification: ZCR helps capture the texture or timbral qualities of music genres, especially those with percussive elements like rock or electronic music.
  - Audio Classification: Distinguish between speech, music, and environmental sounds by analyzing the ZCR over time.
  - Emotion Recognition: In speech, certain emotions (like anger or excitement) tend to increase the ZCR due to rapid articulatory movements.

Totally, it is a simple and effective feature for analyzing the temporal characteristics of an audio signal. It provides valuable insights into the signal's periodicity and noisiness. However, it is often used in combination with other features (e.g., spectral features) for more robust analysis.

8. Linear Predictive Coding (LPC)

LPC is a widely used method in speech processing and audio analysis for representing a signal using a linear predictive model. It analyzes the spectral properties of a signal by estimating the parameters of an autoregressive model that approximates the original signal. LPC is especially popular for applications such as speech compression, synthesis, and feature extraction.

The primary idea of LPC is to model each sample of the signal as a linear combination of past samples. By doing so, LPC efficiently encodes the signal while preserving its essential characteristics, such as formants (resonances in speech signals).

The LPC model represents the signal x(n) at time n as:

$$x(n) = -\sum_{k=1}^{p} a_k \cdot x(n-k) + e(n)$$

Where:

- p: The order of the LPC model (number of coefficients).
- a_k: LPC coefficients (linear prediction coefficients) that represent the model parameters.
- $x(n-k)$: Past pp samples of the signal.
- e(n): The residual error (difference between the predicted and actual values).

The goal is to minimize the residual error e(n) to find the optimal set of LPC coefficients a_k.

- **Purpose:** Models the vocal tract and predicts future samples of a speech signal based on past samples.
- **Process:**
    - **Pre-emphasis:** The audio signal is pre-emphasized by applying a high-pass filter to boost high-frequency components and reduce noise.
    - **Framing and Windowing:** The signal is divided into short overlapping frames to ensure stationarity within each frame. A window function is applied to minimize spectral leakage.
    - **Autocorrelation Analysis:** The autocorrelation method computes the correlation of the signal with its lagged version to estimate the relationship between past samples.
    - **Solve for LPC Coefficients:** The Levinson-Durbin recursion algorithm is used to efficiently solve the linear equations for the LPC coefficients a_k.
    - **Residual Computation:** The residual error e(n) is computed, which can be used for further analysis (e.g., speech synthesis).
- **Applications:**
    - **Speech Compression:** LPC is used in audio codecs (e.g., GSM, VoIP) to compress speech signals by transmitting only the LPC coefficients and the residual error.
    - **Speech Synthesis:** LPC is widely used in text-to-speech (TTS) systems to synthesize speech signals by modeling the vocal tract and using residual excitation.
    - **Speaker Identification:** LPC coefficients capture the unique spectral characteristics of a speaker's voice, making them useful for speaker recognition.
    - **Formant Analysis:** LPC effectively identifies formants (peaks in the spectral envelope), which are important in speech and phoneme analysis.
    - **Music Signal Analysis:** LPC is also used in musical applications to model instruments or remove unwanted components from audio signals.

In summary, LPC is a powerful tool in speech and audio analysis that represents a signal using a small set of linear prediction coefficients. While it assumes linearity and may lose fine-grained spectral details, LPC remains one of the most efficient methods for speech signal modeling.

9. Perceptual Linear Prediction (PLP)

PLP is an advanced feature extraction technique in speech processing that improves upon Linear Predictive Coding (LPC) by incorporating perceptual models of human hearing. PLP is widely used for speech analysis, speaker recognition, and automatic speech recognition systems.

PLP modifies the standard LPC process by incorporating aspects of human auditory perception. This makes it more robust and effective for speech representation. The goal of PLP is to compress the speech spectrum into a perceptually meaningful representation while reducing unnecessary information.

PLP uses critical-band analysis, equal-loudness pre-emphasis, and intensity-loudness compression to mimic the human auditory system

- **Purpose:** Similar to LPC but incorporates auditory perception principles like critical bands and equal loudness curves.
- **Process:**
    - **Pre-emphasis:** Similar to LPC, the speech signal is pre-emphasized to boost high-frequency components and reduce noise.
    - **Windowing and Framing:** The signal is divided into short overlapping frames (e.g., 20–30 ms) to ensure stationarity.
    - **Power Spectrum Calculation:** Apply an FFT to each frame to compute the power spectrum of the signal.
    - **Critical-Band Filtering:** The power spectrum is filtered into a series of critical bands using a bark-scale filter bank. The Bark scale models the frequency resolution of the human ear, grouping nearby frequencies into bands that correspond to auditory critical bands.
    - **Equal-Loudness Pre-Emphasis:** The critical-band spectrum is weighted by an equal-loudness curve, which reflects the fact that human hearing is more sensitive to certain frequency ranges (e.g., mid-range frequencies).
    - **Intensity-Loudness Compression:** The weighted spectrum undergoes a non-linear compression (usually a cubic root function) to model the relationship between sound intensity and perceived loudness.
    - **Linear Prediction:** Perform Linear Predictive Analysis on the compressed spectrum to estimate the spectral envelope, similar to LPC but based on the perceptually processed spectrum.

- **Applications:**
    - **Speaker Recognition:** Captures perceptually relevant speech characteristics, which are useful for identifying individual speakers.
    - **Automatic Speech Recognition:** PLP features are commonly used in speech recognition systems to improve accuracy and robustness, especially in noisy environments.
    - **Emotion Detection:** PLP features can be used in systems that classify emotions based on speech, as they focus on critical aspects of the speech spectrum.
    - **Language Identification:** Helps capture language-specific acoustic properties, improving language classification systems.

Table1: Comparison of PLP vs LPC

Feature	LPC	PLP
Basis	Linear prediction of the signal	Linear prediction with auditory modeling
Spectrum Analysis	No perceptual adjustments	Incorporates human auditory perception
Noise Robustness	Sensitive to noise	More robust to noise
Frequency Scale	Linear	Bark scale (critical bands)
Applications	Basic speech processing	Advanced speech recognition and analysis

To sum up, PLP is a feature extraction technique that enhances the traditional LPC method by incorporating human auditory models. By mimicking the way humans perceive sound, PLP provides a more robust and perceptually meaningful representation of speech. However, its computational complexity and limited frequency resolution must be considered for specific applications.

Similarity Learning

Similarity learning is a branch of machine learning that focuses on training models to recognize the similarity or dissimilarity between data points. It matters because it enables machines to understand patterns, relationships, and structures within data, which is crucial for tasks like recommendation systems, image recognition, and anomaly detection.

While traditional supervised learning focuses on predicting labels based on input data and unsupervised learning aims to find hidden structures within data, similarity learning is somewhat in between.

At its core, Similarity learning is about determining how alike or different two data points are. Imagine you have two photos, and you want to know if they are of the same person. Instead of looking at every pixel, Similarity Learning algorithms will identify key features (like the shape of the eyes or the curve of the mouth) and compare them.

Why is it used? In the vast sea of data, finding patterns or relationships is like finding a needle in a haystack. Similarity learning acts as a magnet, pulling out relevant needles based on their likeness to a given sample.

These algorithms often operate in feature spaces, which are mathematical spaces where data points are represented as vectors. feature vectors. The final layer computes the similarity between these vectors. This method is Ideal for tasks where labeled training examples of dissimilar pairs are hard to find, like signature verification or face verification. A limitation is that they require a lot of data and computational power, plus they might be overkill for simpler tasks where traditional methods could suffice.

In voice similarity detection, similarity learning helps in tasks like speaker verification by ensuring voice embeddings for the same speaker are closer than embeddings for different speakers. A common architecture involves feature extraction The "distance" between these vectors indicates how similar the data points are. The smaller the distance, the more similar they are.

Similarity Learning Methods

Similarity learning hinges on the methods used to measure how alike or different data points are. These methods, often mathematical, provide the foundation for various applications. Let's delve deeper into some of the most common methods:

- **Cosine similarity**

Cosine similarity measures the cosine of the angle between two non-zero vectors. If the vectors are identical, the cosine is 1, indicating maximum similarity. If they are orthogonal (meaning they share no commonality), the cosine is 0. It's particularly suitable for high-dimensional spaces, like text analysis. For instance, in document clustering or when comparing two sets of words to determine their similarity. A limitation of cosine similarity is that it only considers the direction of the vectors, not their magnitude. This means it might not capture the full essence of similarity in cases where magnitude matters.

- **Euclidean distance**

This measures the "straight line" distance between two points in a space. The closer the points are, the more similar they are considered. It's widely used in image recognition and when data can be naturally represented in 2D or 3D spaces. Though it's intuitive and easy to understand, a limitation of Euclidean distance is that In high-dimensional spaces, the concept of "distance" can become less intuitive. Also, all features are given equal importance, which might not always be desired.

- **Siamese networks**

Siamese networks is a neural network approach where two identical subnetworks are defined. These subnetworks take in two inputs and transform them into two via deep networks and subsequent similarity calculation using metrics like cosine similarity

- **Triplet loss**

Used in deep learning, triplet loss involves three data points: an anchor, a positive example (similar to the anchor), and a negative example (different from the anchor). The goal is to ensure that the anchor is closer to the positive example than the negative one by some margin. This method is effective when there's a need to differentiate between very similar-looking data points, like distinguishing between two very similar images of different people. However, it requires careful selection of triplets, especially the negative examples, to ensure effective training. Also, like Siamese networks, it demands substantial data and computational resources.

Understanding the nuances of these methods is crucial. The right method can significantly enhance the accuracy and efficiency of a Similarity Learning task, while the wrong one can lead to subpar results.

Similarity Learning Challenges

While similarity learning offers a plethora of benefits and has revolutionized many sectors, it's not without its challenges.

- **Scalability**: As the volume of data grows exponentially, comparing each data point with every other becomes computationally expensive and time-consuming. This challenge is especially pronounced in real-time applications where quick decisions are essential.
- **Feature selection**: The success of a similarity learning algorithm often hinges on the features chosen for comparison. Not all features are equally important, and identifying the most relevant ones is crucial. Incorrect or redundant features can lead to misleading similarity measures.
- **Noise in data**: Data is rarely perfect. It often contains noise or irrelevant information, which can distort similarity measures. Cleaning and preprocessing data to remove such noise is a significant challenge, especially in large datasets.
- **Overfitting**: This is a common challenge in many machine learning tasks. If a similarity learning model is too complex, it might perform exceptionally well on the training data by memorizing it but fail to generalize to new, unseen data. Striking a balance between model complexity and its generalization capability is crucial to preventing overfitting.
- **Dimensionality**: High-dimensional data can make similarity measures less intuitive and more computationally demanding. Techniques like dimensionality reduction are often required, but they come with the risk of losing important information.

Different Loss Functions used for Similarity learning

Contrastive Loss

Contrastive loss operates on pairs of samples and aims to minimize the distance between embeddings of similar pairs while pushing apart embeddings of dissimilar pairs.

1. **Feature Encoding:** For a pair of inputs, the model computes embeddings
2. **Distance Metric:** Compute the distance between the embeddings, usually using Euclidean distance:
3. **Loss Function:** The loss is computed as:

- Binary label (1 for similar, 0 for dissimilar).
- Distance between the embeddings of the -th pair.
- Margin that defines the minimum allowable distance between dissimilar pairs.

- For similar pairs (), the loss minimizes the distance, pulling the embeddings closer.
- For dissimilar pairs (), the loss encourages to be greater than, pushing the embeddings apart.

Triplet Loss

Triplet loss uses three samples—an anchor x_a, a positive sample x_p(same class as the anchor), and a negative sample x_n(different class).

1. **Feature Encoding:** Compute embeddings for the anchor, positive, and negative samples: f(x_a),f(x_p),f(x_n)
2. **Distance Metrics:**
   - Positive distance:
   $$D_{ap} = \left\| f(x_a) - f(x_p) \right\|_2$$
   - Negative distance:
   $$D_{an} = \| f(x_a) - f(x_n) \|_2$$
3. **Loss Function:** The triplet loss is:
   $$L = \frac{1}{N} \sum_{i=1}^{N} \max(0. D_{ap} - D_{an} + m)$$
   - mm: Margin that enforces a minimum difference between the positive and negative distances.

- The loss minimizes the positive distance Dap (pulling the anchor and positive closer) while ensuring that the negative distance Dan is greater by at least mm (pushing the anchor and negative apart).

Binary Cross-Entropy (BCE)

Binary cross-entropy is often used in similarity learning to classify pairs of samples as similar or dissimilar.

1. **Feature Encoding:** Use a neural network to extract embeddings and for two input samples
2. **Similarity Metric:** Compute the similarity score, often using cosine similarity
3. **Prediction:** Convert the similarity score to a probability using the sigmoid function.

4. **Loss Function:** Compute BCE loss to penalize incorrect predictions.
    - Binary label (1 for similar, 0 for dissimilar).

- BCE directly optimizes the similarity scores to align with binary labels.
- A fixed threshold (e.g., 0.5) can classify pairs during inference.

Table2: Comparison of the Three Losses

Aspect	Contrastive Loss	Triplet Loss	Binary Cross-Entropy
Input Structure	Pairs of samples (x1,x2)	Triplets (xa,xp,xn)	Pairs of samples (x1,x2)
Key Objective	Minimize distance for similar pairs, and maximize distance for dissimilar pairs.	Enforce a margin between positive and negative distances.	Classify pairs as similar or dissimilar.
Advantages	Simpler; easier to compute.	More precise as it directly enforces relative relationships.	Easy to implement; aligns with classification tasks.
Challenges	Requires a well-defined margin mm.	Requires mining hard triplets for effective training.	Sensitive to class imbalance or noisy labels.

All three methods are effective in learning embeddings for tasks like face recognition, speaker verification, and image retrieval. The choice of loss function often depends on the specific requirements of the application and the available computational resources.

References:

- https://www.researchgate.net/publication/371867240_A_Review_of_Recent_Machine_Learning_Approaches_for_Voice_Authentication_Systems?enrichId=rgreq-ecc84ae0bc32a5dab83d80da4c9deeab-XXX&enrichSource=Y292ZXJQYWdlOzM3MTg2NzI0MDtBUzoxMTQzMTI4MTI0Nzk1NDQ3NUAxNzE2OTg4NDMxNzc1&el=1_x_2&_esc=publicationCoverPdf
- https://www.iproov.com/blog/disadvantages-vulnerabilities-voice-biometrics
- https://www.outsourceaccelerator.com/articles/voice-authentication/
- https://www.geeksforgeeks.org/preprocessing-the-audio-dataset/

- https://www.researchgate.net/publication/360336790_An_overview_of_Automatic_Speech_Recognition_Preprocessing_Techniques
- https://www.geeksforgeeks.org/mel-frequency-cepstral-coefficients-mfcc-for-speech-recognition/
- https://www.datacamp.com/blog/what-is-similarity-learning
- https://bigblue.academy/en/similarity-learning
- https://towardsdatascience.com/contrastive-loss-explaned-159f2d4a87ec
- https://towardsdatascience.com/triplet-loss-advanced-intro-49a07b7d8905